


LINEAR REGRESSION (A BEST-FIT LINE PERSPECTIVE)

Inzamam Rahaman

REGRESSION

- Suppose that we have dataset $\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)\}$ where $x_i \in \mathbb{R}^n$ and $y_i \in \mathbb{R}$
- x_i are our independent variable and y_i is our dependent
- Regression is the problem of learning a mapping (model) $f: \mathbb{R}^n \rightarrow \mathbb{R}$ such that $f(x_i) \approx y_i$

REGRESSION

- Suppose that we have dataset $\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)\}$ where $x_i \in \mathbb{R}^n$ and $y_i \in \mathbb{R}$
- Regression is the problem of learning a mapping (model) $f: \mathbb{R}^n \rightarrow \mathbb{R}$ such that $f(x_i) \approx y_i$

mapping is a “generally good” mapping

REGRESSION

- Suppose that we have dataset

$$\mathcal{D} = \{(x_1, y_1), (x_2, y_2), \dots, (x_m, y_m)\} \text{ where } x_i \in \mathbb{R}^n \text{ and } y_i \in \mathbb{R}$$

- Regression is the problem of learning a mapping (model)

$$f: \mathbb{R}^n \rightarrow \mathbb{R} \text{ such that } f(x_i) \approx y_i$$

mapping is a “generally good” mapping

What does it mean
for a mapping to be
“generally good”?

LINEAR REGRESSION

- Space of mappings is infinite!
- Need to devise a structure for our mappings
- Linear regression: restrict our search to a linear function of \mathbb{R}^n

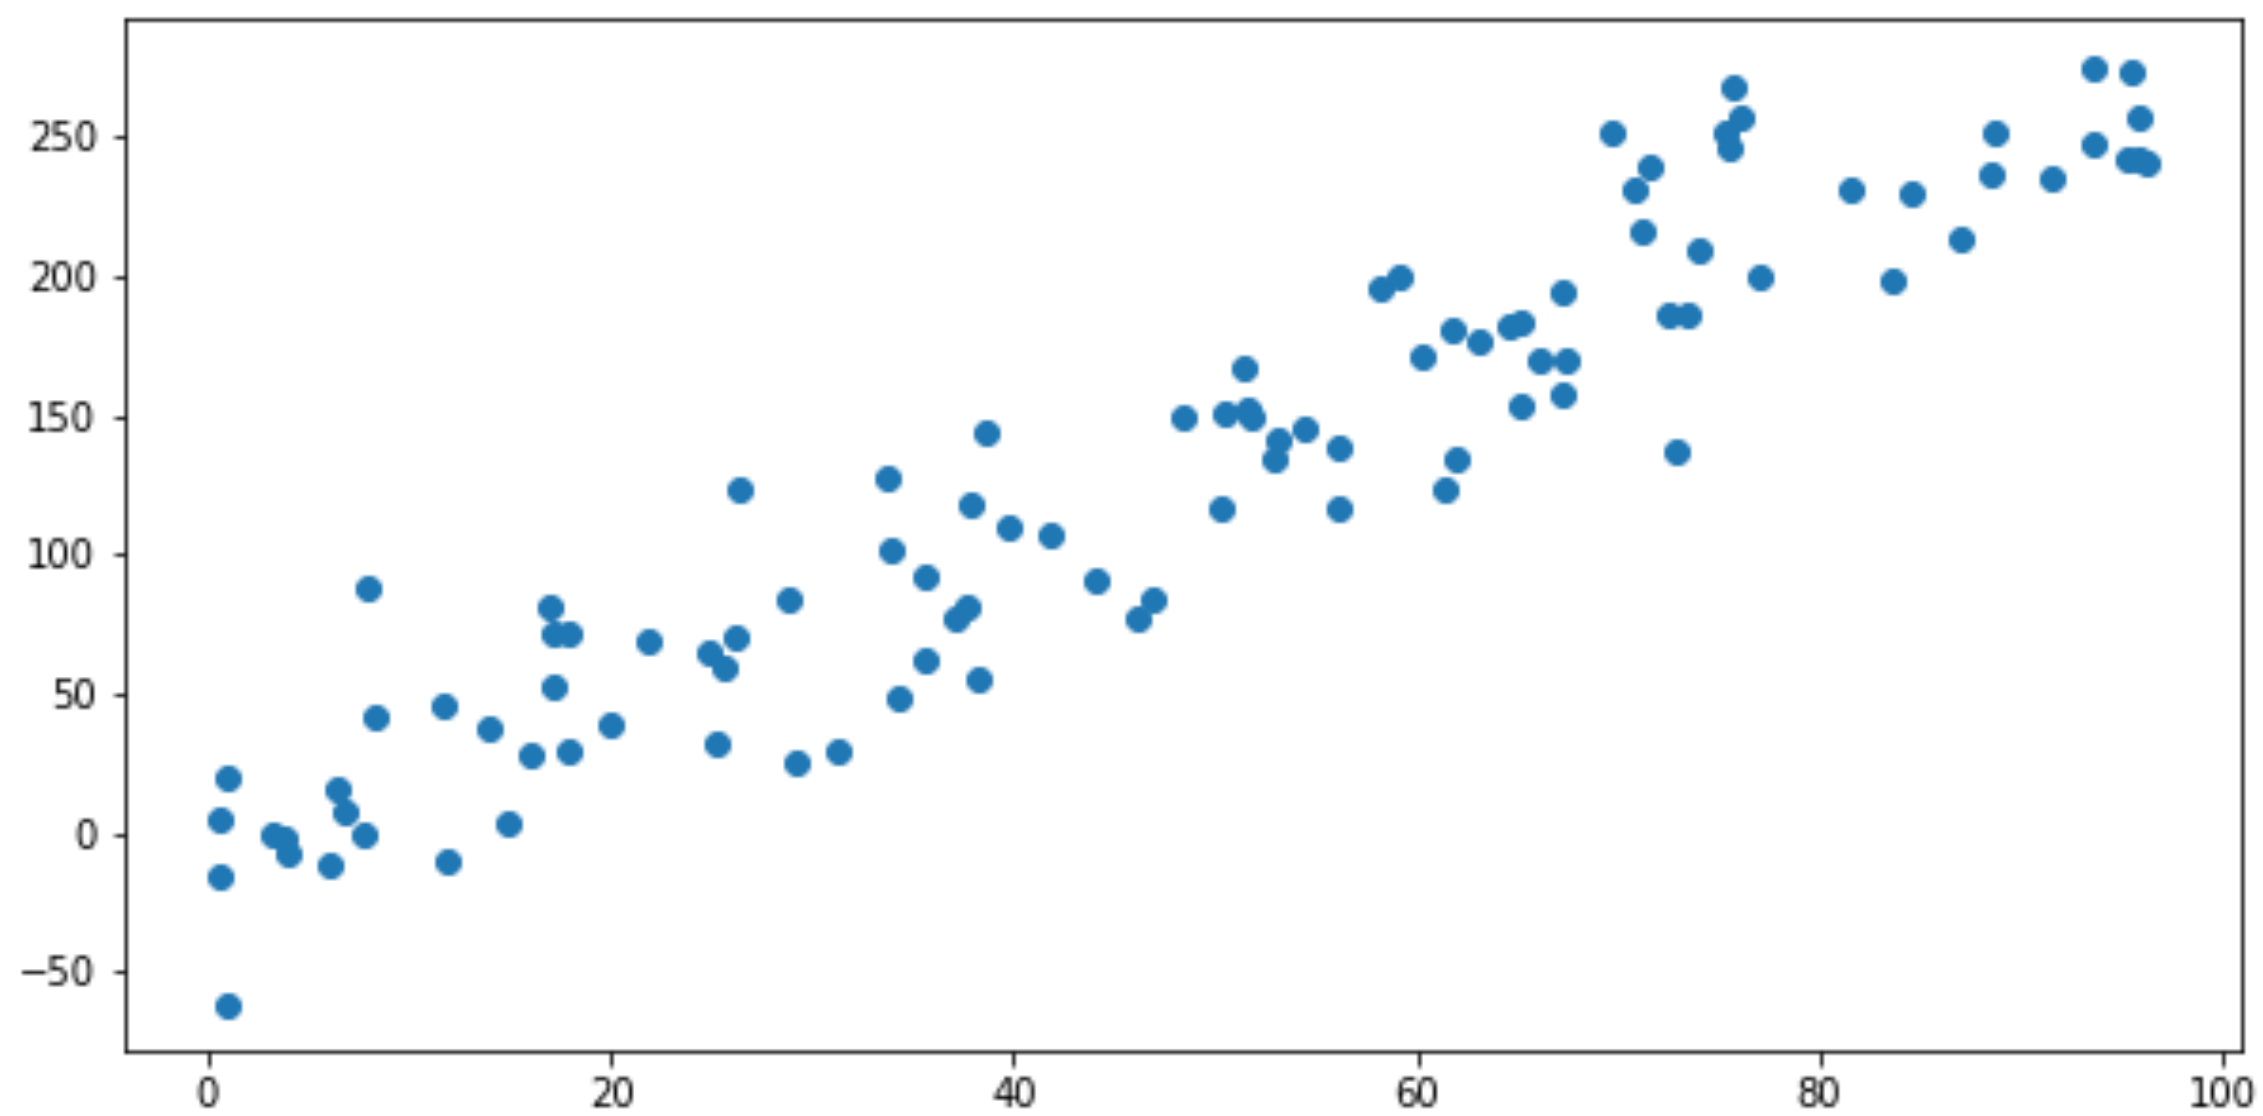



LINEAR REGRESSION


- Linear functions are well understood and malleable
- Assume straight line relationship between input and output
 - think $y = mx + c$
- Applies limitations on model (LR is high-bias low-variance machine learning model family)


LINEAR REGRESSION

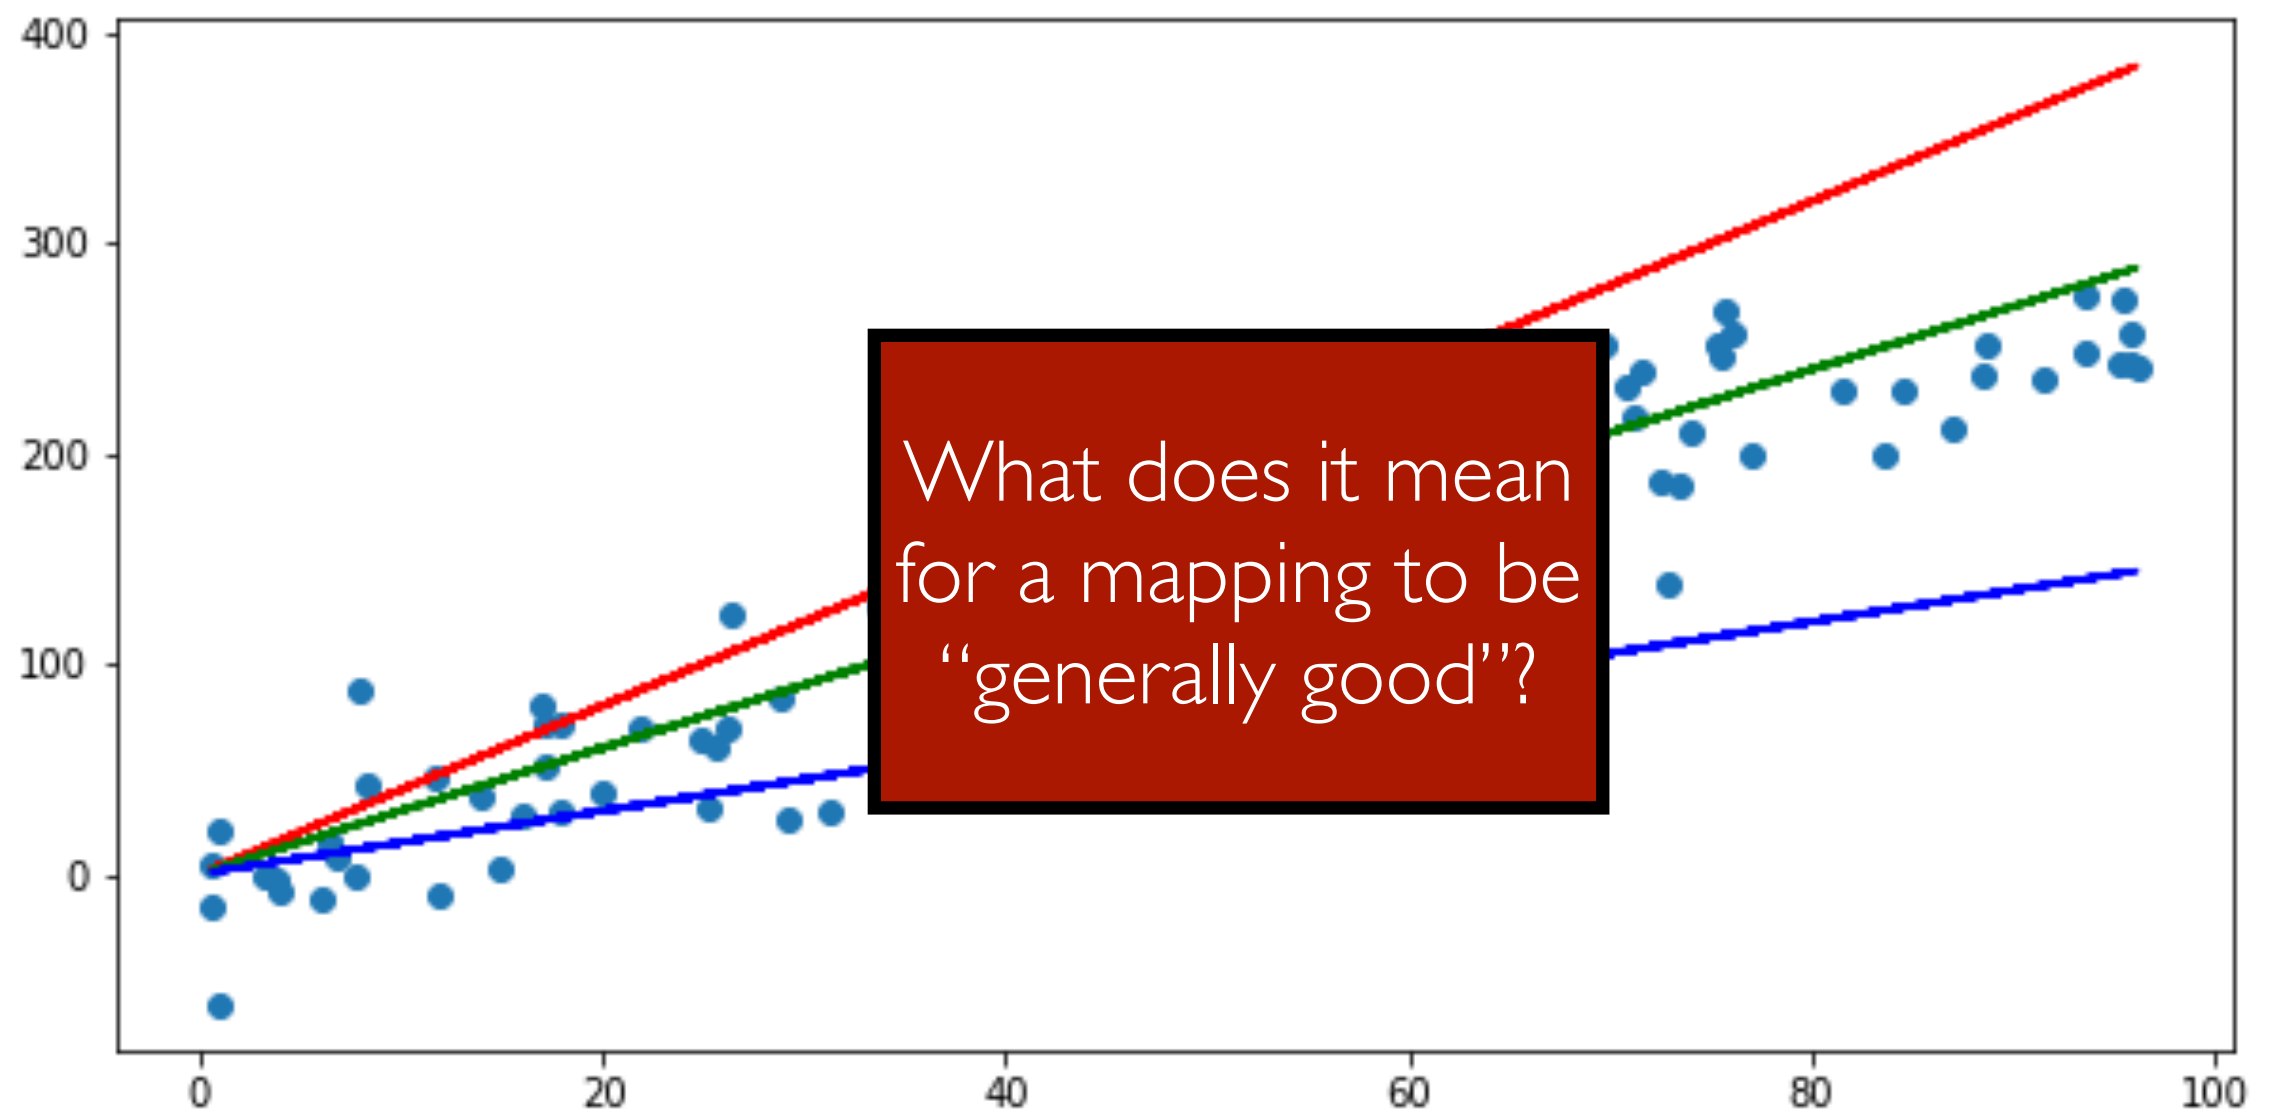
- Because linear, maths scales well to many dimensions
- Will focus on the case of 2 dimensions and assume $c = 0$
 - But maths is trivially extensible











“All models are wrong, but some models are useful”


—George Box

OPTIMISATION AND MACHINE LEARNING

- Machine Learning uses optimisation heavily
- Restrict model space. Make parameters of the model space the inputs to a minimisation problem
 - Want to minimise how “bad” the model performs on **seen** data
 - See Empirical Risk Minimisation for further details

MSE LOSS


- Need to define a loss function to frame the problem of linear regression as a minimisation problem
- Idea: use difference between predicted output and actual output




MSE LOSS

- Let $f_w(x_i) = \hat{y}_i = wx_i$
- Error for example i is $y_i - \hat{y}_i$
- Generally good can be taken to mean a low average error

- $\frac{1}{n} \sum_{i=1}^n y_i - \hat{y}_i ?$





MSE LOSS

- Summing might raw errors might cause errors to cancel out, thereby “drowning” out signal
 - This is undesirable
- Need to add only positive values. Two possibilities:
 - Take the absolute difference (robust regression): $|y_i - \hat{y}_i|$
 - Take the square difference (linear regression): $(y_i - \hat{y}_i)^2$

LINEAR REGRESSION

- For model f_w , where w is our parameter, our loss can be derived as follows:

$$\begin{aligned}\mathcal{L}(D, f_w) &= \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \\ &= \frac{1}{n} \sum_{i=1}^n (y_i - wx_i)^2\end{aligned}$$

Need to solve

$$\begin{array}{l} \text{minimize } \mathcal{L}(D, f_w) \\ w \in \mathbb{R} \end{array}$$

HOW TO SOLVE

- Gradient Descent!

- Recall that $f_3 = f_1 + f_2 \implies \frac{df_3}{dx} = \frac{df_1}{dx} + \frac{df_2}{dx}$

- Also recall the chain rule of differentiation

$$\text{Let } \mathcal{L}_i(D, f_w) = (y - wx_i)^2$$

$$\text{Hence } \frac{d\mathcal{L}_i(D, f_w)}{dw} = -2x_i(y - wx_i)$$

$$\text{Hence } \frac{d\mathcal{L}(D, f_w)}{dw} = \frac{1}{n} \sum_{i=1}^n -2x_i(y - wx_i)$$

$$\text{Let } \mathcal{L}_i(D, f_w) = (y - wx_i)^2$$

$$\text{Hence } \frac{d\mathcal{L}_i(D, f_w)}{dw} = 2x_i(y - wx_i)$$

$$\text{Hence } \frac{d\mathcal{L}(D, f_w)}{dw} = \frac{-2}{n} \sum_{i=1}^n x_i(y - wx_i)$$

Our gradient in gradient descent



WHAT ABOUT THE INTERCEPT

- Can use partial derivative to find update rule for intercept

$$\text{Hence } \frac{d\mathcal{L}(D, f_w)}{dc} = 1$$

IMPROVEMENTS

- In practice, we don't use vanilla gradient descent
- Use a variations such as stochastic gradient descent and friends (AdaGrad, RMSProp, etc...)
- In SGD, we don't calculate loss over all data points per iteration, only a sample
 - SGD is faster in practice, but may not find as a good a point as GD

REGULARIZATION

- To prevent overfitting, we use regularization
- Penalize “size” of weights using l_1 or l_2 norm